

Suyash Tambe 22070126117 AIML B2

```

1 import torch
2 import torch.nn as nn
3 import torch.optim as optim
4 import torchvision
5 import torchvision.transforms as transforms
6 from torch.utils.data import DataLoader
7 import matplotlib.pyplot as plt

1 transform = transforms.Compose([
2     transforms.Resize((64, 64)),
3     transforms.ToTensor(),
4     transforms.Normalize((0.5,), (0.5,))
5 ])
6
7
8 dataset = torchvision.datasets.CIFAR10(root="./data", train=True, download=True, transform=transform)
9 dataloader = DataLoader(dataset, batch_size=64, shuffle=True)
10
11 device = "cuda" if torch.cuda.is_available() else "cpu"
12

```

 Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz  
 100% |██████████| 170M/170M [00:12<00:00, 13.3MB/s]  
 Extracting ./data/cifar-10-python.tar.gz to ./data

```

1 # Define Autoencoder (AE)
2 class Autoencoder(nn.Module):
3     def __init__(self, latent_dim=128):
4         super().__init__()
5         self.encoder = nn.Sequential(
6             nn.Linear(64*64*3, 512),
7             nn.ReLU(),
8             nn.Linear(512, latent_dim)
9         )
10        self.decoder = nn.Sequential(
11            nn.Linear(latent_dim, 512),
12            nn.ReLU(),
13            nn.Linear(512, 64*64*3),
14            nn.Tanh()
15        )
16
17    def forward(self, x):
18        x = x.view(x.size(0), -1)
19        encoded = self.encoder(x)
20        decoded = self.decoder(encoded)
21        return decoded.view(x.size(0), 3, 64, 64)
22

```

```

1 # Define Variational Autoencoder (VAE)
2 class VariationalAutoencoder(nn.Module):
3     def __init__(self, latent_dim=128):
4         super().__init__()
5         self.encoder = nn.Sequential(
6             nn.Linear(64*64*3, 512),
7             nn.ReLU()
8         )
9         self.fc_mu = nn.Linear(512, latent_dim)
10        self.fc_logvar = nn.Linear(512, latent_dim)
11        self.decoder = nn.Sequential(
12            nn.Linear(latent_dim, 512),
13            nn.ReLU(),
14            nn.Linear(512, 64*64*3),
15            nn.Tanh()
16        )
17
18    def reparameterize(self, mu, logvar):
19        std = torch.exp(0.5 * logvar)
20        eps = torch.randn_like(std)
21        return mu + eps * std
22
23    def forward(self, x):
24        x = x.view(x.size(0), -1)
25        x = self.encoder(x)

```

```

26     mu, logvar = self.fc_mu(x), self.fc_logvar(x)
27     z = self.reparameterize(mu, logvar)
28     decoded = self.decoder(z)
29     return decoded.view(x.size(0), 3, 64, 64), mu, logvar

1 # Train Autoencoder
2 def train_autoencoder(model, dataloader, epochs=10, lr=1e-3):
3     optimizer = optim.Adam(model.parameters(), lr=lr)
4     criterion = nn.MSELoss()
5     model.to(device)
6
7     for epoch in range(epochs):
8         for images, _ in dataloader:
9             images = images.to(device)
10            outputs = model(images)
11            loss = criterion(outputs, images)
12            optimizer.zero_grad()
13            loss.backward()
14            optimizer.step()
15            print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")
16    return model

1 # Train Variational Autoencoder
2 def train_vae(model, dataloader, epochs=10, lr=1e-3):
3     optimizer = optim.Adam(model.parameters(), lr=lr)
4     model.to(device)
5
6     def vae_loss(recon_x, x, mu, logvar):
7         recon_loss = nn.MSELoss()(recon_x, x)
8         kl_divergence = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
9         return recon_loss + kl_divergence
10
11    for epoch in range(epochs):
12        for images, _ in dataloader:
13            images = images.to(device)
14            recon_images, mu, logvar = model(images)
15            loss = vae_loss(recon_images, images, mu, logvar)
16            optimizer.zero_grad()
17            loss.backward()
18            optimizer.step()
19            print(f"Epoch [{epoch+1}/{epochs}], Loss: {loss.item():.4f}")
20    return model
21

1
2 # Visualization Function
3 def visualize_reconstruction(model, dataloader, is_vae=False):
4     model.eval()
5     images, _ = next(iter(dataloader))
6     images = images.to(device)
7     with torch.no_grad():
8         if is_vae:
9             reconstructed, _, _ = model(images)
10        else:
11            reconstructed = model(images)
12
13    fig, axes = plt.subplots(2, 8, figsize=(10, 4))
14    for i in range(8):
15        axes[0, i].imshow(images[i].cpu().permute(1, 2, 0) * 0.5 + 0.5)
16        axes[0, i].axis('off')
17        axes[1, i].imshow(reconstructed[i].cpu().permute(1, 2, 0) * 0.5 + 0.5)
18        axes[1, i].axis('off')
19    plt.show()
20
21 # Train and visualize Autoencoder
22 ae = Autoencoder(latent_dim=128)
23 ae = train_autoencoder(ae, dataloader)
24 print("Visualizing Autoencoder Reconstruction")
25 visualize_reconstruction(ae, dataloader)
26
27 # Train and visualize Variational Autoencoder
28 vae = VariationalAutoencoder(latent_dim=128)
29 vae = train_vae(vae, dataloader)
30 print("Visualizing Variational Autoencoder Reconstruction")
31 visualize_reconstruction(vae, dataloader, is_vae=True)
32

```

Epoch [1/10], Loss: 0.0280  
Epoch [2/10], Loss: 0.0197  
Epoch [3/10], Loss: 0.0217  
Epoch [4/10], Loss: 0.0251  
Epoch [5/10], Loss: 0.0214  
Epoch [6/10], Loss: 0.0206  
Epoch [7/10], Loss: 0.0177  
Epoch [8/10], Loss: 0.0174  
Epoch [9/10], Loss: 0.0147  
Epoch [10/10], Loss: 0.0190

Visualizing Autoencoder Reconstruction



Epoch [1/10], Loss: 6.5653  
Epoch [2/10], Loss: 1.3779  
Epoch [3/10], Loss: 0.3987  
Epoch [4/10], Loss: nan  
Epoch [5/10], Loss: nan  
Epoch [6/10], Loss: nan  
Epoch [7/10], Loss: nan  
Epoch [8/10], Loss: nan  
Epoch [9/10], Loss: nan  
Epoch [10/10], Loss: nan

Visualizing Variational Autoencoder Reconstruction

